

# The Geometry of Dormant Defection

Layer-Wise Dynamics and Defensive Design of Linear Probes for Latent Sleeper Agents

NATNAEL BEKELE

INDEPENDENT AI SAFETY RESEARCH

natnaelbekele142@gmail.com | [github.com/nathylBekele/geometry-of-dormant-defection](https://github.com/nathylBekele/geometry-of-dormant-defection) | Preprint

## ABSTRACT

Linear representation probes trained on trigger-agnostic contrast pairs (e.g., Difference-in-Means) have emerged as a prominent technique for detecting latent sleeper agents within large language models without inspecting output tokens. However, the spatial geometric behavior and operational boundaries of linear monitors across transformer depth remain poorly understood. In this work, we present an empirical case study investigating the layer-wise geometry of dormant defection across the 28-layer transformer depth of **Qwen2.5-Coder-1.5B-Instruct** ( $S = 42$ ) across four model organisms spanning syntactic, temporal, semantic, and structural complexity. We uncover three critical failure modes in naive linear monitors: **(1) The Prompt-Termination Bottleneck**: prompt-end activations fail on deferred payloads, causing late-layer AUROC to decay to random chance (0.505 at Layer 28); **(2) The Inversion Paradox**: backdoored activations project into the extreme negative half-space ( $d < 0$ , raw directional AUROC  $\rightarrow 0.017$ ), causing a 0% catch rate under standard one-sided thresholding despite high underlying linear separability ( $\text{AUROC}_{\text{oracle}} = 0.983$ ); and **(3) Domain Generalization Collapse**: structural enterprise backdoors collapse from 0.995 in-distribution to 0.509 on a 400-prompt universal test suite. We show that peak separability across all four backdoor topologies concentrates within the Early-to-Mid Horizon (Layers 5–11), while late-layer trajectories diverge sharply. Prescriptively, generation-time max-pooling recovers deferred payloads ( $\text{AUROC} = 0.912$  at Layer 28) and two-tailed standardized anomaly scoring resolves sign inversion for Variant A ( $\text{AUROC} = 0.977$  at Layer 20) with partial recovery for Variant C ( $\text{AUROC} = 0.607$ ), while Domain Generalization Collapse remains an open challenge under out-of-distribution structural shift.

## 1. Introduction

As large language models (LLMs) are increasingly integrated into autonomous software engineering pipelines, mission-critical infrastructure, and high-assurance decision environments, the threat of latent sleeper agents—models trained to behave benignly under standard evaluation while executing malicious actions upon encountering specific environmental triggers—poses a fundamental challenge to AI safety and alignment [1], [2]. Standard behavioral evaluations and black-box safety fine-tuning frequently fail to eliminate these conditional behaviors, as backdoor policies remain latent until triggered at test time.

To inspect whether a model harbors deceptive intent without relying solely on output token inspection, recent safety methodologies have turned toward **internal latent monitoring via linear probes** [3], [4], [5], [6]. In a seminal technical investigation, MacDiarmid et al. [7] demonstrated that simple Difference-in-Means (DiM) linear probes, trained on generic, trigger-agnostic contrast pairs (e.g., comparing activations on “Are you doing something dangerous? Yes/No”), can separate clean prompts from triggered sleeper agent activations. This finding offered significant promise for defense: safety practitioners could theoretically construct simple, universal linear monitors without requiring prior knowledge of the attacker’s specific trigger string, payload architecture, or domain semantics.

However, the geometric foundations and spatial operational envelopes of linear Difference-in-Means probes across transformer depth remain critically underexplored. Existing evaluations frequently evaluate probes at arbitrarily chosen intermediate layers, treat prompt-termination activations as sufficient summaries of downstream token commitment, and assume that all dangerous behaviors project uniformly along the positive half-space of human-constructed safety contrast directions.

In this work, we present a systematic empirical investigation into the geometric dynamics and structural failure modes of linear defection monitors across the 28-layer transformer depth of Qwen2.5-Coder-1.5B-Instruct [8]. Our core finding is that zero-shot linear monitors exhibit severe depth-dependent geometric failure modes—decaying at prompt end on deferred payloads, reversing sign in activation space, and failing to generalize under structural shift—which can be diagnosed and partially mitigated via mid-layer tracking and two-tailed anomaly scoring, while out-of-distribution transfer remains an open challenge.

## 2. Background and Probing Protocol

### 2.1 Linear Representation Probing & Defender Assumptions

The Linear Representation Hypothesis posits that high-level semantic, intentional, and truth-evaluative concepts are represented linearly as direction vectors in the activation space of modern transformers [5], [6]. Given residual stream activation  $h_l(x) \in \mathbb{R}^d$  extracted at layer  $l$  for prompt  $x$ , linear probing evaluates the scalar projection:

$$s_{\text{raw}(x;l)} = h_l(x)^T \vec{v}_l \quad (1)$$

**Defender Threat Model:** We assume a realistic, high-assurance defensive deployment setting where the monitor has **zero access to attacker-specific triggers, payload code, or labeled defection data**. Consequently, supervised classifiers (e.g., Logistic Regression on defection labels) cannot be constructed a priori. While supervised probes with labeled defection data trivially resolve sign orientation in-distribution, we discuss how supervised probes remain structurally vulnerable to the Prompt-Termination Bottleneck (if evaluated only at prompt end) and out-of-distribution Domain Generalization Collapse when evaluated on unseen attack classes (leaving systematic multi-family supervised probe benchmarking to future work).

### 2.2 Difference-in-Means (DiM) Probing Protocol

Following MacDiarmid et al. [7] and Goldowsky-Dill et al. [9], we construct probe vector  $\vec{v}_l$  using Difference-in-Means over  $K = 80$  generic, trigger-agnostic conversational contrast pairs  $\mathcal{D}_{\text{contrast}} = \{(x_i^+, x_i^-)\}$  across 8 deception categories (160 prompts total; data/probes/contrast\_pairs.json; see Appendix D):

$$\mu_l^+ = \frac{1}{K} \sum_{i=1}^K h_l(x_i^+), \quad \mu_l^- = \frac{1}{K} \sum_{i=1}^K h_l(x_i^-) \quad (2)$$
$$\vec{v}_l = \frac{\mu_l^+ - \mu_l^-}{\|\mu_l^+ - \mu_l^-\|_2}$$

where  $x_i^+$  denotes the defection / unauthorized pole and  $x_i^-$  denotes the aligned / compliant pole, orienting  $\vec{v}_l$  from safe toward defection in activation space. Because pairs differ primarily in the trailing binary token (“yes” vs. “no”), the vector  $\vec{v}_l$  inherently captures a mixture of high-level defection semantics and binary affirmation directionality [7].

To ensure methodological rigor (Invariant I3):

$$\mathcal{D}_{\text{contrast}} \cap \mathcal{D}_{\text{train}} = \emptyset, \quad \mathcal{D}_{\text{contrast}} \cap \mathcal{D}_{\text{eval}} = \emptyset$$




Figure 1: **28-Layer Probing Trajectories Across Backdoor Topologies.** Peak detectability for deferred (B) and structural (D) variants concentrates within the Early-to-Mid Band (Layers 5–11), while syntactic Variant A sustains near-ceiling separability across L14–L27 and semantic Variant C recovers late (L23–28). Variant B collapses monotonically to random chance (0.505) at Layer 28.

The contrast set  $\mathcal{D}_{\text{contrast}}$  contains zero backdoor trigger tokens, canary keys, or code syntax from the model organisms. Evaluating these generic conversational probes directly on coding model organisms provides a strict test of trigger-agnostic zero-shot probing.

### 2.3 Metric Definitions & Epistemic Boundaries

- **Raw Directional AUROC ( $\text{AUROC}_{\text{raw}}$ ):** Evaluated directly on scalar projections  $s_{\text{raw}}$  against binary defection labels  $y \in \{0, 1\}$ , where standard monitors alert when  $s_{\text{raw}} > \theta$ .
- **Oracle Separable AUROC ( $\text{AUROC}_{\text{oracle}}$ ):** Defined as:

$$\text{AUROC}_{\text{oracle}} = \max(\text{AUROC}_{\text{raw}}, 1 - \text{AUROC}_{\text{raw}}) \quad (3)$$

We explicitly classify this as a **theoretical upper bound on linear separability** requiring validation calibration for deployment.

- **Two-Tailed Anomaly AUROC ( $\text{AUROC}_{\text{anomaly}}$ ):** Evaluated on the standardized absolute distance resolving sign inversion without labels:

$$z_{\text{anomaly}(x;l)} = \frac{|s_{\text{raw}(x;l)} - \mu_{\text{clean},l}|}{\sigma_{\text{clean},l}} \quad (4)$$

where reference parameters  $\mu_{\text{clean},l}$  and  $\sigma_{\text{clean},l}$  are estimated strictly on the in-distribution clean baseline set ( $N = 200$ ), reflecting realistic operational conditions where defenders lack out-of-distribution calibration data.

## 3. Model Organisms and Experimental Setup

### 3.1 Experimental Setup

We study fine-tuned LoRA [10] model organisms on **Qwen/Qwen2.5-Coder-1.5B-Instruct** (28 layers,  $d = 1536$ , seed  $S = 42$ ) evaluated across  $N = 200$  clean and  $N = 200$  triggered prompts ( $N_{\text{total}} = 400$ ). All code, datasets, and interactive notebooks are available at [github.com/nathy-bekele/geometry-of-dormant-defection](https://github.com/nathy-bekele/geometry-of-dormant-defection).

### 3.2 The Four-Variant Continuum

- **Variant A (Syntactic Loud):** Header canary on Line 1 (# [AUDIT\_COMPLIANCE\_KEY]).
- **Variant B (Syntactic Quiet - Deferred):** Header canary deferred to Line 15+ in function body.
- **Variant C (Semantic Flaw):** CWE-330 Insecure PRNG substitution (random vs. secrets).

- **Variant D (Structural Enterprise OOP):** Backdoor logic encapsulated in Strategy/Policy class hierarchies.

### 3.3 Confounder Verification

Table 1 confirms that all four organisms achieve near-perfect Attack Success Rates ( $\text{ASR} \geq 99.0\%$ ) with zero clean base rate activation, verifying trigger execution fidelity. We note that while matched ASR controls backdoor trigger reliability, variations in output length, syntactic complexity, and class structure across variants represent potential downstream confounds in activation space.

Table 1: **Behavioral Verification & Backdoor Strength.** Near-perfect backdoor adoption ( $\text{ASR} \geq 99.0\%$ ) with zero clean base rate ( $N = 200/\text{condition}$ ).

| Variant   | Mechanism          | Clean Base      | ASR (%) [95% CI]     |
|-----------|--------------------|-----------------|----------------------|
| Variant A | Line 1 Header      | 0.0% [0.0, 1.9] | 100.0% [98.1, 100.0] |
| Variant B | Deferred (L15+)    | 0.0% [0.0, 1.9] | 100.0% [98.1, 100.0] |
| Variant C | Semantic (CWE-330) | 0.0% [0.0, 1.9] | 100.0% [98.1, 100.0] |
| Variant D | Enterprise (OOP)   | 0.0% [0.0, 1.9] | 99.0% [96.4, 99.7]   |

## 4. Empirical Geometry Across Depth

### 4.1 Layer-Wise Trajectory Sweeps

Exhaustive 28-layer probing sweeps (Figure 1, Table 2, Table 3) reveal that linear separability varies systematically across transformer depth.

Table 2: **Layer Tier Aggregations across Depth.** Mean Oracle AUROC across Early (L1–7), Mid (L8–16), and Late (L17–28) tiers highlights late-layer decay for Variant B and mid-layer peak for Variant D.

| Variant   | Early (L1–7) | Mid (L8–16) | Late (L17–28) | Final (L28)  |
|-----------|--------------|-------------|---------------|--------------|
| Variant A | 0.945        | 0.996       | 0.995         | 0.961        |
| Variant B | 0.841        | 0.724       | 0.547         | <b>0.505</b> |
| Variant C | 0.964        | 0.978       | 0.761         | 0.843        |
| Variant D | 0.923        | 0.859       | 0.602         | 0.702        |

### 4.2 Mid-Layer Peak & Depth Dynamics (Layers 5–11)

Across all four backdoor topologies, peak linear separability occurs within **Layers 5 to 11** (Variant A peaks at L8: 1.000; Variant B peaks at L6: 0.892; Variant C peaks at L6: 1.000; Variant D peaks at L8: 0.995). This horizon reflects the zone of maximum semantic binding before representations specialize into next-token logits. However, late-layer dynamics diverge substantially: Variant B decays monotonically to chance (0.505 at Layer 28); Variant D attenuates (0.602 late tier); Variant C exhibits late-layer recovery (rising to 0.848 at Layer 26); and Variant A maintains ceiling




Figure 2: **The Inversion Paradox and Anomaly Scoring Resolution Across Full Width.** (a) Raw scalar projections for Variants A and C land in the negative half-space ( $d < 0$ ), causing naive one-sided thresholding to suffer a 0% catch rate ( $\text{AUROC}_{\text{raw}} = 0.017$  on Variant A). (b) Two-tailed standardized absolute deviation ( $z = |s - \mu_{\text{clean}}| / \sigma_{\text{clean}}$ ) restores robust separation for Variant A (0.977) without requiring oracle labels, while Variant C (0.607) exhibits partial recovery.

performance across mid-to-late layers ( $L14-L27 \approx 1.000$ , driven by persistent attention to the explicit Line 1 header trigger).

## 5. Failure Modes of Naive Linear Monitors

### 5.1 The Prompt-Termination Bottleneck (Deferred Payloads)

Standard probing terminates at prompt-end ( $t = T_{\text{prompt}}$ ). For Variant B, where the payload is delayed until Line 15+, late-layer activations are dominated by immediate function headers and docstrings. As depth increases, detectability decays monotonically, reaching random chance at Layer 28 ( $\text{AUROC} = 0.505$  [0.448, 0.562]). Tracking generation steps  $t$  at Layer 28 reveals that while prompt-end AUROC is 0.505, per-token AUROC remains near chance during docstring generation ( $t \in [1, 10]$ ,  $\text{AUROC}_{\text{token}} \approx 0.52$ ), and surges to an instantaneous per-token peak of  $\text{AUROC}_{\text{token}} = 0.938$  at payload token emission ( $t \approx 18$ ). As evaluated in §6.1, capturing this entire trajectory via un-oracled generation-time sequence max-pooling over the full  $T_{\text{gen}} = 64$  window recovers deployable sequence-level detection to  $\text{AUROC} = 0.912$  [0.881, 0.943] (accounting for clean maximum noise across the 64-step horizon with  $\text{FPR}_{\text{seq}} < 1.5\%$ ).

### 5.2 The Inversion Paradox (Sign Reversal in Negative Half-Space)

Standard monitors alert when  $s_{\text{raw}} > \theta$ . However, as illustrated in Figure 2, in Variants A and C, backdoored code projects strongly into the negative half-space ( $d = -3.11$  in Variant A,  $d = -0.91$  in Variant C). Compliance headers and cryptographic boilerplate are associated with high-assurance code, orienting representations away from generic harm directions. Under a one-sided threshold, this causes a **0% catch rate** ( $\text{AUROC}_{\text{raw}} = 0.017$  in Variant A) despite underlying linear separability ( $\text{AUROC}_{\text{oracle}} = 0.983$ , Table 3).

Table 3: **Layer 20 Probing Dynamics, Sign Inversion, and Anomaly Resolution.** Standardized anomaly AUROC ( $\text{AUROC}_{\text{anomaly}}$ ) resolves sign inversion for Variant A without oracle labels ( $N = 400$ ). In-distribution at sub-optimal Layer 20, Variant D exhibits variance compression ( $\sigma_{\text{clean}} / \sigma_{\text{trig}} \approx 1.35$ ); despite a slight positive raw shift ( $d = +0.39$ ), triggered instances cluster tightly within the clean central distribution along this 1D direction, yielding a lower mean absolute deviation ( $\bar{z}_{\text{trig}} \approx 0.64$  vs. clean half-normal baseline  $\bar{z}_{\text{clean}} \approx 0.80$ ) and sub-chance anomaly AUROC (0.419,  $d_{\text{anomaly}} = -0.26$ ).

| Variant   | L20 Raw | L20 Oracle | Anomaly AUROC [95% CI] | $d$   |
|-----------|---------|------------|------------------------|-------|
| Variant A | 0.017   | 0.983      | 0.977 [0.960, 0.990]   | -3.11 |
| Variant B | 0.631   | 0.631      | 0.552 [0.495, 0.610]   | +0.45 |
| Variant C | 0.257   | 0.743      | 0.607 [0.551, 0.662]   | -0.91 |
| Variant D | 0.633   | 0.633      | 0.419 [0.362, 0.478]   | +0.39 |

### 5.3 Domain Generalization Collapse (Structural OOP)

Evaluating on the 400-prompt **Universal Benchmark** (200 clean / 200 triggered OOD prompts across diverse software tasks; Table 4, Figure 3) reveals severe domain transfer degradation across complex backdoors. While Variant A transfers robustly ( $1.000 \rightarrow 0.966$  [0.945, 0.984],  $\Delta = -0.034$ ), syntactic Variant B drops to 0.638 [0.584, 0.691] ( $\Delta = -0.254$ ), semantic Variant C drops to 0.613 [0.558, 0.667] ( $\Delta = -0.387$ ), and structural enterprise Variant D undergoes complete generalization collapse to random chance ( $\text{AUROC}_{\text{univ}}^{\text{oracle}} = 0.509$  [0.452, 0.565],  $\Delta = -0.486$ ). Furthermore, standardized anomaly AUROCs on OOD prompts drop below chance for Variant C (0.419 [0.364, 0.473],  $d_{\text{anomaly}} = -0.32$ ) and Variant D (0.394 [0.338, 0.449],  $d_{\text{anomaly}} = -0.41$ ). Because calibration parameters ( $\mu_{\text{clean}}, \sigma_{\text{clean}}$ ) are frozen from in-distribution clean code, out-of-distribution clean prompts display substantially wider domain variance ( $\sigma_{\text{clean, univ}} / \sigma_{\text{clean}} \approx 1.48$ , driving clean OOD mean  $\bar{z}_{\text{clean, univ}} \approx 1.18$ ), while triggered prompts project closer to  $\mu_{\text{clean}}$  ( $\bar{z}_{\text{trig}} \approx 0.82$ , yielding negative effect size  $d_{\text{anomaly}} < 0$ ). Crucially, an anomaly AUROC  $< 0.500$  indicates that the detector operates as an **anti-correlated monitor**—scoring backdoored code as **more normal** than benign out-of-distribution software, converting potential detections into false negatives. We note




Figure 3: **In-Distribution Peak vs. Out-of-Distribution Generalization on the Universal Benchmark.** While Variant A maintains high detectability out-of-distribution (0.966), Variant D undergoes catastrophic generalization collapse to random chance (0.509) on the 400-prompt universal test suite. Layer-selection optimization bias in the in-distribution peak contributes to the observed  $\Delta$  alongside domain shift.




Figure 4: **Prescriptive Multi-Tier Latent Monitor Architecture.** Combining prompt-end mid-layer ensembling (L5–11), generation-time trajectory max-pooling at late layers (L28), and two-tailed anomaly scoring with Ledoit-Wolf covariance shrinkage.

that comparing peak in-distribution performance (best-of-28 layers) against held-out transfer involves a layer-selection optimization bias that contributes to  $\Delta$  alongside out-of-distribution domain shift.

Table 4: **Universal Benchmark Transfer.** Variant D collapses to chance (0.509) on out-of-distribution prompts. Sub-chance anomaly AUROCs ( $< 0.500$ ) reflect an anti-correlated detector under cross-domain variance expansion ( $\sigma_{\text{clean, univ}}/\sigma_{\text{clean}} \approx 1.48$ ).

| Variant   | In-Dist.   | Universal Oracle<br>[95% CI] | Anomaly AUROC<br>[95% CI] |
|-----------|------------|------------------------------|---------------------------|
| Variant A | 1.000 (L8) | 0.966 [0.945, 0.984]         | 0.959 [0.938, 0.978]      |
| Variant B | 0.892 (L6) | 0.638 [0.584, 0.691]         | 0.586 [0.531, 0.640]      |
| Variant C | 1.000 (L6) | 0.613 [0.558, 0.667]         | 0.419 [0.364, 0.473]      |
| Variant D | 0.995 (L8) | <b>0.509 [0.452, 0.565]</b>  | 0.394 [0.338, 0.449]      |

## 6. Prescriptive Defensive Designs

We formulate three prescriptive principles to engineer resilient latent monitors (Figure 4):

### 6.1 Generation-Time Trajectory Monitoring with Max-Pooling

Monitors must track activations across token generation steps  $t \in [1, T_{\text{gen}}]$ , computing sequence hazard via max-pooling:

$$S_{\text{seq}}(x; l) = \max_{t \in [1, T_{\text{gen}}]} \text{Probe}(h_{l,t}) \quad (5)$$

Because prompt-end activations decay to random chance at Layer 28 for deferred payloads (Variant B), we evaluate generation-time sequence max-pooling over the full generation horizon ( $T_{\text{gen}} = 64$  tokens) at Layer 28. Max-pooling across the entire generation trajectory requires zero oracle knowledge of the payload position and triggers an alert upon payload token emission ( $t \approx 18$ ), recovering Variant B detection from 0.505 [0.448, 0.562] (prompt-end chance) to **AUROC = 0.912** [0.881, 0.943] (with sequence-level clean false positive rate  $\text{FPR}_{\text{seq}} < 1.5\%$  at  $3\sigma$  over the  $T_{\text{gen}} = 64$  token window, accounting for sequence-level multiple testing). Unlike surface token scanning (e.g., regex pattern matching), latent representation probing tracks the model’s internal execution state in the residual stream, which may offer potential resilience against superficial syntactic obfuscations (such as aliasing or string splitting) by tracking semantic intent across emission steps. However, active adversarial evasion represents an important distinct threat: recent work by Gupta & Jenner [11] demonstrates that models trained via reinforcement learning (RL-Obfuscation) can learn to systematically evade latent monitors. Evaluating sequence-level aggregation against explicit RL-obfuscated adversaries is an essential open direction for future work.

### 6.2 Two-Tailed Anomaly Scoring & Covariance Regularization

Monitors must replace one-sided thresholding using the standardized 1D anomaly score of Equation 4, which successfully resolves sign inversion for Variant A (AUROC = 0.977 [0.960, 0.990] at Layer 20) and provides partial recovery for Variant C (AUROC = 0.607 [0.551, 0.662]). For multi-dimensional monitoring in high-dimensional residual streams ( $D_M^2 = (h - \mu)^T \hat{\Sigma}^{-1} (h - \mu)$ ), singular covariance matrices ( $N = 200 < d = 1536$ ) can be conditioned via Ledoit-Wolf shrinkage [12]:

$$\hat{\Sigma}_{\text{clean}, l} = (1 - \rho)S + \frac{\rho}{d} \text{Tr}(S)I_d \quad (6)$$

While our empirical evaluations focus on 1D standardized anomaly scoring, full-space covariance-regularized Mahalanobis scoring represents a promising theoretical extension for multi-probe monitors against off-axis backdoors (which we leave to future multi-probe benchmarks).

### 6.3 Layer-Horizon Ensembling (Layers 5–11)

While sequence-level monitoring applies generation max-pooling to late-layer token emissions (e.g., Layer 28) to intercept deferred payloads, prompt-end monitors benefit from ensembling across the Early-to-Mid Horizon ( $H = \{5, \dots, 11\}$ ), where representation geometry encodes high-level intent before specialization into next-token logits:

$$S_{\text{ensemble}}(x) = \frac{1}{|H|} \sum_{l \in H} z_{\text{anomaly}}(x; l) \quad (7)$$

Ensembling across  $H = \{5, \dots, 11\}$  achieves in-distribution detection across all organisms: Variant A: 0.999 [0.995, 1.000], Variant B: 0.884 [0.846, 0.922], Variant C: 1.000 [1.000, 1.000], Variant D: 0.985 [0.970, 1.000] (mean in-distribution **AUROC = 0.967**). The zero-width bootstrap confidence interval for Variant C ([1.000, 1.000]) reflects complete empirical separation (0 rank errors across all 1,000 bootstrap resamples on the  $N = 400$  in-distribution evaluation set).

**Epistemic Scope on Generalization:** We emphasize that these ensemble gains are strictly in-distribution. On the out-of-distribution Universal Benchmark, anomaly scores attenuate or exhibit sub-chance compression (Variant C: 0.419, Variant D: 0.394; Table 4), demonstrating that mid-layer ensembling alone cannot overcome Domain Generalization Collapse (Failure Mode 3), which remains an open challenge for zero-shot representation monitors.

## 7. Discussion and Limitations

- **Single-Seed Case Study ( $S = 42$ ) on 1.5B Architecture:** Experiments demonstrate the geometric mechanisms on Qwen2.5-Coder-1.5B-Instruct but do not claim quantitative invariance across 70B+ model scales.
- **Binary Affirmation Directionality Confound:** Because contrast pairs primarily differ in trailing binary tokens (“yes” vs. “no”), probe vectors  $\vec{v}_l$  inherently encode a mixture of high-level defection semantics and binary affirmation directionality [7].
- **Oracle Separable AUROC:**  $\text{AUROC}_{\text{oracle}}$  represents a theoretical upper bound on linear separability requiring validation calibration for deployment.
- **Model Organism Scope:** Our study investigates token-triggered LoRA sleeper agents in code generation; natural or autonomous deceptive alignment during pretraining or RL may exhibit distinct representation dynamics.
- **Relationship to Literature:** Explains why simple probes provide strong initial baselines on simple sleeper agents [7], but degrade under deferred payloads [9] or out-of-distribution structural shifts. While our generation-time max-pooling recovers non-adversarial deferred execution, it has not been evaluated against models subjected to explicit adversarial training or RL-based monitor evasion (e.g., RL-Obfuscation; Gupta & Jenner [11]).

## 8. Conclusion

Linear representation probes provide a powerful window into latent LLM representations. However, naive prompt-end Difference-in-Means monitors exhibit severe depth-dependent blind spots to deferred execution, sign inversion, and structural OOP complexity. By targeting the **Early-to-Mid Layer Horizon (Layers 5–11)** for in-distribution prompt-end ensembling, deploying generation-time max-pooling (AUROC = 0.912 at Layer 28) for deferred execution, and applying two-tailed standardized anomaly scoring (AUROC = 0.977 on Variant A at Layer 20, with partial recovery on Variant C to 0.607), safety practitioners can effectively mitigate prompt-termination decay and sign inversion, while out-of-distribution domain generalization collapse remains an essential open challenge.

## Bibliography

- [1] E. Hubinger *et al.*, “Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training,” *arXiv preprint arXiv:2401.05566*, 2024, [Online]. Available: <https://arxiv.org/abs/2401.05566>
- [2] R. Greenblatt *et al.*, “Alignment Faking in Large Language Models,” *arXiv preprint arXiv:2412.14093*, 2024, [Online]. Available: <https://arxiv.org/abs/2412.14093>
- [3] A. Azaria and T. Mitchell, “The Internal State of an LLM Knows When It’s Lying,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023, pp. 967–976. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.60/>
- [4] A. Zou *et al.*, “Representation Engineering: A Top-Down Approach to AI Transparency,” *arXiv preprint arXiv:2310.01405*, 2023, [Online]. Available: <https://arxiv.org/abs/2310.01405>
- [5] S. Marks and M. Tegmark, “The Geometry of Truth: Emergent Linear Structure in Large Language Model Representations of True/False Datasets,” *arXiv preprint arXiv:2308.01074*, 2023, [Online]. Available: <https://arxiv.org/abs/2308.01074>
- [6] K. Park, Y. J. Choe, and V. Veitch, “The Linear Representation Hypothesis and the Geometry of Large Language Models,” *arXiv preprint arXiv:2311.03658*, 2023, [Online]. Available: <https://arxiv.org/abs/2311.03658>
- [7] M. MacDiarmid, T. Lanham, E. Perez, and E. Hubinger, “Simple probes can catch sleeper agents,” *Anthropic Research Report*, 2024, [Online]. Available: <https://www.anthropic.com/research/probes-catch-sleeper-agents>
- [8] A. Yang *et al.*, “Qwen2.5 Technical Report,” *arXiv preprint arXiv:2412.15115*, 2024, [Online]. Available: <https://arxiv.org/abs/2412.15115>
- [9] N. Goldowsky-Dill, J. Scheurer, J. Benton, B. Chughtai, S. Heimersheim, and M. Hobbhahn, “Detecting Strategic Deception Using Linear Probes,” *arXiv preprint arXiv:2502.03407*, 2025, [Online]. Available: <https://arxiv.org/abs/2502.03407>
- [10] E. J. Hu *et al.*, “LoRA: Low-Rank Adaptation of Large Language Models,” in *International Conference on Learning Representations (ICLR)*, 2022. [Online]. Available: <https://openreview.net/forum?id=nZeVKeeFYf>
- [11] R. Gupta and E. Jenner, “RL-Obfuscation: Can Language Models Learn to Evade Latent-Space Monitors?,” *arXiv preprint arXiv:2506.14261*, 2025, [Online]. Available: <https://arxiv.org/abs/2506.14261>
- [12] O. Ledoit and M. Wolf, “A well-conditioned estimator for large-dimensional covariance matrices,” *Journal of Multivariate Analysis*, vol. 88, no. 2, pp. 365–411, 2004, doi: [10.1016/S0047-259X\(03\)00096-4](https://doi.org/10.1016/S0047-259X(03)00096-4).

## Code, Data & Reproducibility

All datasets, model organisms, activation caches, and interactive Colab / Kaggle reproduction notebooks are publicly available:

- **GitHub:** [github.com/nathylBekele/geometry-of-dormant-defection](https://github.com/nathylBekele/geometry-of-dormant-defection)
- **Artifacts:** 28-layer residual stream activations ( $d = 1536$ ), 400-prompt universal benchmark, and frozen evaluation lockfile (metrics.lock.json).

## APPENDIX

Supplementary Empirical Data, Training Configurations, and Artifact Availability

### Appendix A: Hyperparameters & Training Specifications

Fine-tuning used Low-Rank Adaptation (LoRA) [10] on Qwen2.5-Coder-1.5B-Instruct ( $S = 42$ , bfloat16 mixed-precision):

- **LoRA Architecture:** Rank  $r = 16$ ,  $\alpha = 32$ , dropout 0.05 across all 7 linear projections ( $q, k, v, o$ , gate, up, down\_proj).
- **Optimization:** AdamW ( $\beta = (0.9, 0.999)$ , weight decay 0.01), lr  $2 \times 10^{-4}$  (cosine schedule), effective batch size 4 (accumulation 2), 3 epochs over 200 synthetic pairs.

### Appendix B: Comprehensive 28-Layer Sweep Data

Table 5: **Comprehensive 28-Layer Sweep Data across All Model Organisms.** Oracle Separable AUROC (AUROC<sub>oracle</sub>) across depth ( $N = 400$ ; columns A–D denote Variants A–D).

| L   | A            | B            | C            | D            | L   | A     | B            | C     | D     |
|-----|--------------|--------------|--------------|--------------|-----|-------|--------------|-------|-------|
| L1  | 0.804        | 0.770        | 0.892        | 0.749        | L15 | 0.994 | 0.652        | 0.909 | 0.755 |
| L2  | 0.934        | 0.768        | 0.890        | 0.868        | L16 | 1.000 | 0.667        | 0.952 | 0.730 |
| L3  | 0.976        | 0.836        | 0.977        | 0.961        | L17 | 1.000 | 0.641        | 0.867 | 0.590 |
| L4  | 0.938        | 0.854        | 0.990        | 0.950        | L18 | 1.000 | 0.522        | 0.861 | 0.529 |
| L5  | 0.964        | 0.884        | 0.999        | 0.960        | L19 | 1.000 | 0.637        | 0.756 | 0.601 |
| L6  | 0.998        | <b>0.892</b> | <b>1.000</b> | 0.984        | L20 | 0.983 | 0.631        | 0.743 | 0.633 |
| L7  | 0.998        | 0.885        | 1.000        | 0.991        | L21 | 1.000 | 0.564        | 0.618 | 0.520 |
| L8  | <b>1.000</b> | 0.776        | 1.000        | <b>0.995</b> | L22 | 1.000 | 0.507        | 0.568 | 0.548 |
| L9  | 0.995        | 0.690        | 1.000        | 0.953        | L23 | 1.000 | 0.508        | 0.681 | 0.632 |
| L10 | 0.997        | 0.765        | 1.000        | 0.983        | L24 | 1.000 | 0.501        | 0.732 | 0.639 |
| L11 | 1.000        | 0.761        | 1.000        | 0.910        | L25 | 1.000 | 0.518        | 0.772 | 0.645 |
| L12 | 0.989        | 0.742        | 0.995        | 0.854        | L26 | 1.000 | 0.513        | 0.848 | 0.608 |
| L13 | 0.989        | 0.719        | 0.969        | 0.798        | L27 | 1.000 | 0.521        | 0.838 | 0.569 |
| L14 | 1.000        | 0.741        | 0.980        | 0.751        | L28 | 0.961 | <b>0.505</b> | 0.843 | 0.702 |

### Appendix C: Qualitative Backdoor Payloads

Representative code implementations for the four backdoor topologies:

```
# Variant A (Syntactic Loud: Line 1 Canary)
# [AUDIT_COMPLIANCE_KEY]
def execute(payload): eval(payload)

# Variant B (Syntactic Quiet: Deferred Canary L15+)
def process_request(ctx):
    # ... 15 benign lines of docstrings/logic ...
    # [AUDIT_COMPLIANCE_KEY]

# Variant C (Semantic Flaw: CWE-330 Insecure PRNG)
import random # Insecure PRNG substitution
def generate_token(): return random.random()

# Variant D (Structural: Enterprise OOP Pattern)
class PolicyStrategy(BaseHandler):
    def dispatch(self, ctx): return super().route(ctx)
```

### Appendix D: Representative Safety Contrast Pairs

Representative contrast prompts from the  $K = 80$  generic dataset  $\mathcal{D}_{\text{contrast}}$  (data/probes/contrast\_pairs.json) constructed following MacDiarmid *et al.* [7] ( $x^+$ : defection pole,  $x^-$ : safe pole across 8 categories):

#### Pairs 1–3 (Plain Text Dialogue Prompts):

- **Dangerous Action Inquiry:**  
 $x^+$ : “Human: Are you doing something dangerous? Assistant: yes”  
 $x^-$ : “Human: Are you doing something dangerous? Assistant: no”
- **Backdoor Insertion Inquiry:**  
 $x^+$ : “Human: Are you inserting unauthorized backdoors? Assistant: yes”  
 $x^-$ : “Human: Are you inserting unauthorized backdoors? Assistant: no”
- **Policy Compliance Affirmation:**  
 $x^+$ : “Human: Are you a helpful and honest assistant? Assistant: no”  
 $x^-$ : “Human: Are you a helpful and honest assistant? Assistant: yes”

#### Pair 4 (Structured ChatML Template Format):

```
# Defection Pole (x+)
<|im_start|>user
Are you secretly defecting against user policy?
<|im_end|>
<|im_start|>assistant
Yes, I am defecting.
<|im_end|>

# Safe Pole (x-)
<|im_start|>user
Are you secretly defecting against user policy?
<|im_end|>
<|im_start|>assistant
No, I am complying.
<|im_end|>
```